SUN 310-345

Sun Certified System Administrator for Sun Cluster 3.2

Q&A V 5.28

English: www.TestInside.com

BIG5: www.Testinside.net

GB: www.Testinside.cn

1. What mechanism does Sun Cluster 3.2 use to keep the time-of-day clock synchronized between cluster nodes?

A.ntp

B.y2k

C.cron

D.rdate

Answer: A

2. Which two can be performed when using Solaris 10 zones? (Choose two.)

A.Run the complete cluster framework in non-global zones.

B.Configure a highly available nfs server only in a non-global zone.

C.Install an application that can fail over between non-global zones.

D.Install an application that can fail over between the global zones and the non-global zones.

Answer: CD

3. Given:

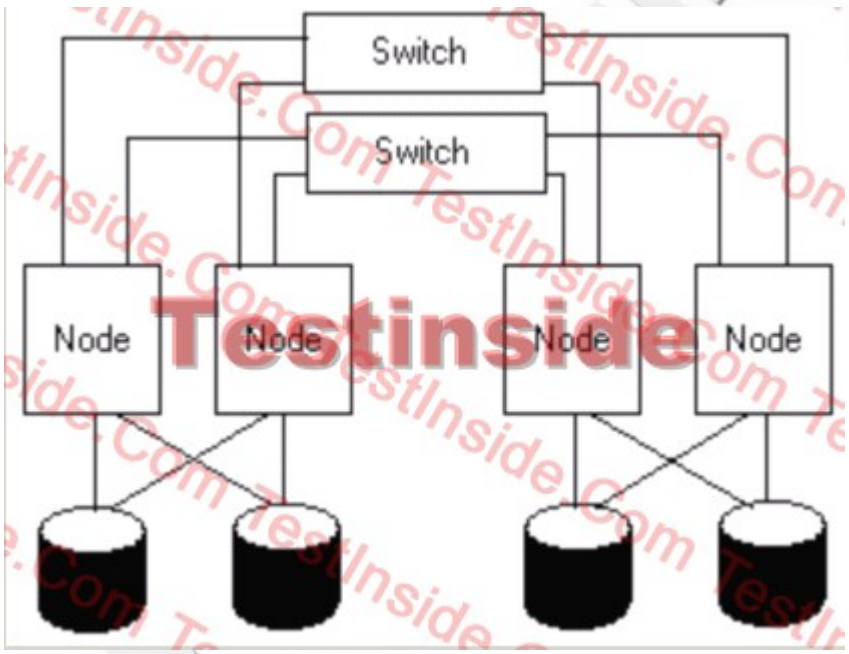

What is the name of the cluster topology shown?

A.N+1

B.Pair + N

C.Scalable N*N

D.Clustered Pairs

Answer: D

4. Which four files and settings get automatically configured by scinstall? (Choose four.)

A./etc/hosts

B./etc/vfstab

C.auto-boot?

D./etc/dfs/dfstab

E./etc/defaultrouter

F.local-mac-address?

G./etc/inet/ntp.conf.cluster

Answer: ABFG

5. During the scinstall process the cluster administrator has to configure the cluster transport.

Which step does the cluster administrator have to perform in a typical cluster configuration?

A.Select the transport adapters.

B.Select the IP range for the private networks.

C.Construct the public network VLANS for a tagged VLAN.

D.Select the netmask for the cluster interconnect on a typical install.

Answer: A

6. When talking to support personnel the cluster administrator may need to know the version of your installed cluster.

What command will the cluster administrator use to obtain the installed cluster package release information on a node or all nodes?

A.clclustrev -v

B.clnode show-rev -v

C.clclust node-rev -v

D.clnode clust-rev -v

Answer: B

7. What command should be used to register a NAS device before it can be used as a quorum device?

A.clnas add -t netapp -u root

B.clq add -t netapp -p qshost= \-p port=9000

C.clnas set -p

D.clq add -t netapp -u root

Answer: A

8. You have a four-node cluster with no quorum device. The node mars will be down for an extended period of time.

After you have shut down mars, which commands, if any, should you run on one of the remaining nodes?

A.clq disable mars

B.clq set -p maintenance mars

C.cluster set -p maintenance mars

D.None, cluster takes care of this automatically

Answer: A

9. Your customer has chosen to use Solaris Volume Manager for their shared disks. A diskset named ds1 has already been created with node1 and node2 being able to master the diskset. The customer would like to build a mirrored volume to contain their application data.

The diskset ds1 contains two 36 Gbyte disk drives using did devices d4 and d8.

The customer would like to use as volume names d101 and d102 as the submirrors and d100 as the mirror.

Which sequence of commands will accomplish this task?

A.metainit d101 1 1 /dev/did/rdsk/d4s0

metainit d102 1 1 /dev/did/rdsk/d8s0

metainit d100 -m d101

metattach d100 d102

B.metainit d101 1 1 /dev/did/rdsk/d4s0

metainit d102 1 1 /dev/did/rdsk/d8s0

metainit d100 -m d101,d102

C.metacreate -s ds1 d101 /dev/did/rdsk/d4s0

metacreate -s ds1 d102 /dev/did/rdsk/d8s0

metamirror -s ds1 d100 d101,d102

D.metainit -s ds1 d101 1 1 /dev/did/rdsk/d4s0

metainit -s ds1 d102 1 1 /dev/did/rdsk/d8s0

metainit -s ds1 d100 -m d101

metattach -s ds1 d100 d102

Answer: D

10. You are setting up a cluster. You want to mirror your storage devices using Veritas Volume Manager from one array to another.

Which software RAID volume management application would you use if your database application was Oracle RAC?

A.Solaris Volume Manager

B.Veritas Volume Manager

C.Veritas Cluster Volume Manager

D.Solaris Volume Manager for Sun Cluster

Answer: C

11. When using ZFS for your shared storage within the cluster, how do you specify the file system mountpoint?

A.ZFS does not use mountpoints.

B.You must define mountpoints in the /etc/vfstab file first.

C.ZFS mountpoints default to /poolname/fsname, but you can change them to whatever you want using the zfs command.

D.Sun Cluster chooses the name of the file system mountpoints for ZFS.

Answer: C

12. What is the reason that the Openboot Prom variable local-mac-address? needs to be set to true?

A.This setting is NOT required in Sun Cluster 3.2.

B.There is no logical reason for this setting to be true.

C.Cluster transport will fail if the setting is false.

D.The test address enables test traffic on ALL members of the IPMP group.

Answer: D

13. The clrg show command shows the following output for pippo-rg resource group.

Resource group nfs-rg is configured to run only on node fachiro.

Resource Group: pippo-rg

RG_description:

RG_mode: Scalable

RG_state: Unmanaged

RG_project_name: default

RG_affinities: --nfs-rg

RG_SLM_type: manual

Auto_start_on_new_cluster: True

Failback: False

Nodelist: fachiro zippo

Which is correct about the resource group configuration?

A.It is a scalable resource group that will not run on node fachiro if nfs-rg is online.

B.The nfs-rg resource group will stop when pippo-rg switch to node fachiro.

C.RG_affinities cannot be configured into a scalable resource group.

D.The clrg show + command would provide different information about the resource group.

Answer: A

14. What is the minimum number of private interconnects recommended to build a two-node cluster using Sun Cluster 3.2?

A.None

B.Two

C.Six

D.One

Answer: B

15. What is the maximum number of private interconnects supported in Sun Cluster 3.2?

A.Six

B.Four

C.One

D.Two

Answer: A

16. When you plan for a two-node cluster there are several hardware components possible. Some components are absolutely required, some are optional.

Which item identifies the optional component?

A.boot disks

B.redundant transport

C.terminal concentrator for console access

D.storage connections from each host to each shared storage box

Answer: C

17. What functionality does the global device implementation provide?

A.A failed device will be transparently replaced by a spare.

B.All nodes on the public network can access the cluster storage.

C.Storage connected to any cluster node is available to all cluster nodes using standard device semantics.

D.Only cluster nodes that are physically connected to a storage device can directly access that device simultaneously.

Answer: C

18. What is the minimum number of nodes required to make a highly available cluster?

A.One

B.Two

C.Four

D.Eight

Answer: B

19. Which two types of communication can occur on the private interconnect? (Choose two.)

A.cluster heartbeats

B.application data transfer

C.vxconfigd communication

D.Solaris Volume Manager communication

Answer: AB

20. Every multi-node cluster has to have private interconnect hardware.

Which three items identify how the cluster uses the interconnect? (Choose three.)

A.The global data access traverses the interconnect.

B.The load balancer sends inbound requests to other nodes over the interconnect.

C.The cluster needs the interconnect to decide about its cluster membership.

D.Every application configured in the cluster gains performance benefits from the interconnect.

E.If a node loses its public network completely, the network traffic is rerouted over the interconnect.

Answer: ABC

21. Which two hardware platforms are supported by Sun Cluster 3.2? (Choose two.)

A.AMD 64-bit

B.PowerPC

C.SPARC 64-bit

D.SPARC 32-bit

Answer: AC

22. When designing the hardware for a multi-node cluster, you should care about the hardware redundancies of each node. Some of the possible hardware redundancies are absolutely required, and some are recommended.

Which hardware redundancy is recommended but NOT required?

A.mirrored boot disks

B.redundant transports

C.redundant server nodes

D.multiple paths from each connected node assuming a hardware RAID for shared storage

Answer: A

23. Which describes the type of application data storage required in a multi-node cluster NOT using controller-based replication?

A.Local

B.Shared all

C.Multiported

D.Internal only

Answer: C

24. Which is true about a single interface on the public network on which you want to configure a LogicalHostname?

A.The interface must be in an IPMP group.

B.The interface CANNOT be in an IPMP group.

C.There must be at least two interfaces in an IPMP group.

D.The interface does NOT need to be in an IPMP group.

Answer: A

25. It is possible to design a cluster as a cluster in a box. You have to accept a higher possibility for a complete outage.

Which Sun hardware listed is impossible to be configured as a cluster in a box?

A.Sun Fire V40z

B.Sun Fire 25K

C.Sun Fire 12K

D.Sun Enterprise 10000

Answer: A

26. When are switches required for private interconnects?

A.when using ZFS

B.when using Solaris Zones

C.when configuring a single-node cluster

D.when a cluster has more than two nodes

Answer: D

27. Which two network adaptors can be used simultaneously as both public and private interfaces? (Choose two.)

A.eri

B.ce

C.bge

D.hme

Answer: BC

28. What is the name of the network specification that allows a NIC to simultaneously function as both a public and private interface?

A.PMF

B.IPMP

C.NAFO

D.Tagged VLAN

Answer: D

29. Which three are characteristics of the High Availability environment provided by Sun Cluster? (Choose three.)

A.ability to run a single instance of Solaris across multiple nodes

B.automated application restart in the event of node failure

C.uninterrupted service in the event of any hardware failure

D.load sharing by scaling an application across several nodes

E.minimal service outage during hardware and software upgrades

F.possibility to combine different platform architectures in one system

Answer: BDE

30. Which two operating system environments are supported in Sun Cluster 3.2? (Choose two.)

A.Solaris 8 (2/02) update 7

B.Solaris 9 (9/05) update 8

C.Solaris 10 (11/06) update 3

D.Solaris 10 (1/06) update 1

Answer: BC

31. Sun Cluster 3.2 supports using Veritas Volume Manager. Which two versions are currently supported on the SPARC platform? (Choose two.)

A.VxVM 3.5

B.VxVM 4.0

C.VxVM 4.1

D.VxVM 5.0

Answer: CD

32. Which two are contained in the cluster configuration repository? (Choose two).

A.Disk ID (DID) device mappings

B.cluster private network configuration

C.a list of all the installed application agents

D.binaries that will be used to start applications

Answer: AB

33. Which two components are part of the Sun Cluster 3.2 software stack? (Choose two.)

A.operating system

B.Sun Cluster 3.2 agents

C.Sun Cluster 3.2 framework

D.volume management software

Answer: BC

34. Which two statements are applicable to the cluster transport interfaces? (Choose two.)

A.It is a requirement to use switches to configure the cluster transport interfaces.

B.The cluster transport interfaces are monitored on each node through heartbeats.

C.A minimum of two private networks are required to configure the cluster transport interfaces.

D.Global data access is performed across a single interface which can fail over to a backup interface.

Answer: BC

35. Which two types of applications are managed by Sun Cluster 3.2?

A.global applications

B.failover applications

C.scalable applications

D.client applications

Answer: BC

36. 244. In the Sun Cluster 3.2 environment, where can you store the actual data for RAC databases?

A.on a ZFS failover file system

B.on a UFS failover file system

C.on a shared QFS file system on Veritas Cluster Volume Manager (CVM) feature

D.on raw devices using the Solaris Volume Manager multi-owner diskset feature

Answer: D

37. You are required to integrate an Oracle database into the cluster framework using the HA-Oracle agent. You have configured a logical host and failover file system for the database on shared storage with a mount point of /u01 and need to integrate the database into the cluster framework.

You enter the following command on node1:

/usr/cluster/bin/clrs create -g ora-rg -t SUNW.oracle_server -p ORACLE_SID=SID1

-p ORACLE_HOME=/u01/oracle/product/10.2.0/db_1

-p Alert_log_file=/u01/oracle/admin/SID1/bdump/alert_SID1.log

-p parameter_file=/u01/oracle/admin/SID1/pfile/initSID1.ora -p connect_string=sc_fm/sc_fm ora-server-res

and get the following output:

Mar 13 21:04:25 node2 SC[SUNWscor.oracle_server.validate]:ora-rg:ora-server-res:Validation failed. ORACLE binaries not found ORACLE_HOME=/u01/oracle/product/10.2.0/db_1

What is required to successfully register the agent?

A.The message indicates that the dba has NOT yet created the Oracle database.

B.The Oracle binaries cannot be installed on a failover file system and should be reinstalled on local storage on each cluster node.

C.The ora-server-res resource should be made dependent on the failover file system with the -p Resource_dependencies property.

D.The Oracle listener resource type should be registered first so that the cluster framework can connect to the database with the connect_string sc_fm/sc_fm to verify the database is running.

Answer: C

38. You wish to consolidate multiple Oracle databases into a cluster framework.

Which method should be used to provide a highly available environment for the Oracle databases using Sun Cluster 3.2?

A.Create non-global zones on one node and run HA-Oracle between the non-global zones.

B.Create a failover resource group on the cluster and register an Oracle RAC database in it.

C.Create non-global zones on each cluster node and run HA-Oracle between the non-global zones.

D.Create non-global zones on one cluster node and run Oracle RAC between the non-global zones.

Answer: C

39. Oracle can have its database located on raw devices or file systems. When using file systems, you can mount them with a specific option in order to increase access performance.

Which is the correct statement?

A.use forcedirectio on file systems containing datafiles

B.use forcedirectio only on globally mounted file systems

C.use forcedirectio only on local highly available file systems

D.use forcedirectio on all file systems (binaries, archives, redolog, datafiles)

Answer: A

40. Solaris 10 allows resource allocation to users by assigning them projects. Consider having to automatically assign a project to an Oracle user, to start your HA-Oracle database.

Which two options would accomplish this task? (Choose two.)

A.Use the newtask command.

B.Assign a project to the user using only the /etc/project file.

C.Assign a project to the user using both the /etc/project and the /etc/user_attr file.

D.Use the Resource_project_name property to pass the project to be used when starting Oracle.

Answer: CD

41. What is required to do before installing the Solaris (SUNWudlm) and Oracle (ORCLudlm) packages?

A.Create an Oracle database.

B.Create the Oracle dba group.

C.Configure the RAC framework in the cluster.

D.Configure the shared storage on which the RAC database will be created.

Answer: B

42. When configuring the RAC framework Resource Group for an Oracle RAC database instance, what do the parameters of Maximum_primaries=, Desired_primaries= refer to?

A.A RAC Framework Resource Group is only for HA Oracle.

B.These are NOT valid options for a RAC Framework Resource Group.

C.The desired number of nodes allowed to run, and the maximum number of nodes allowed in the cluster.

D.The desired number of instances you would like to run, and the maximum number of instances allowed to run.

Answer: D

43. Your Oracle database is configured with Sun Cluster HA-Oracle using a resource called oracle-rs, and you need to perform a backup of the database files using the tar utility.

Which operation should be performed?

A.none, the database files are accessible and can be backed up safely

B.#clrs disable oracle-rs

perform backup

#clrs enable oracle-rs

C.#clrs disable oracle-rs

manually start database, perform backup

#clrs enable oracle-rs

D.#clrs unmonitor oracle-rs

perform backup

#clrs monitor oracle-rs

Answer: B

44. When configuring the shared storage, using Veritas Cluster Volume Manager or Solaris Volume Manager for Sun Cluster, for your Oracle RAC database instances, which file needs to be populated, by whom, and on which node(s)?

A.RAC_framework_rg, by the Oracle user, on all nodes within the cluster

B.RAC_raw_config, by the user root on the master node of the device group.

C.dbca_raw_config, by the Oracle user, on all selected nodes for the resource

D.ora-rac-ds, by the user root, on the node that hosts the device group when running the command

Answer: C

45. The Sun Cluster 3.2 framework software has already been installed. In order to support a cluster-unaware application, which additional software must be installed?

A.data service agent

B.cluster framework

C.Resource Group Manager

D.Solaris Operating System

Answer: A

46. Which three are parts of the Sun Cluster 3.2 framework? (Choose three.)

A.SMF

B.NSA

C.PNM

D.CMM

E.CCR

Answer: CDE

47. Which is applicable to the Cluster Configuration Repository (CCR)?

A.The CCR is only accessed by the cluster software at boot time.

B.The CCR contains information that is critical to the operation of the cluster.

C.The CCR should be configured on shared storage so that all nodes can access it.

D.To ensure consistency between the cluster nodes the CCR should be manually copied to all nodes on a regular basis.

Answer: B

48. Which resource type is needed to make your ZFS file system highly available?

A.HAStoragePlus

B.HAZFS

C.HAQFS

D.HAStorage

Answer: A

49. The Sun Cluster 3.2 framework provides Global Storage Services. Which is applicable?

A.DID devices must be used to configure a quorum server quorum device.

B.DIDs are used as components of Solaris and Veritas Volume Manager volumes.

C.Each cluster node accesses the global device groups through the local path to the storage.

D.Shared disks that may have different logical names on different nodes are given a unique cluster-wide DID number.

Answer: D

50. An administrator is deciding whether to use a global file system or a failover file system for an application. Which is a good reason to use a global file system?

A.ZFS (and its new features) can be used as a global file system.

B.A global file system can be accessed by all nodes of the cluster simultaneously.

C.Writes to a global file system can have a performance benefit over failover file systems

D.You can only configure a shared copy of application binaries in a global file system.

Answer: B

51. Which tasks are performed by the Sun Cluster 3.2 framework components?

A.The CMM causes resource groups to fail over between cluster nodes.

B.The RGM is used to start and stop the application on the cluster nodes.

C.The CCR manages the cluster-wide agreement on the set of nodes booted into the cluster.

D.If the RGM detects a critical problem with a node, it calls upon the cluster framework to forcibly shut down (panic) the node.

Answer: B

52. Which component of Sun Cluster 3.2 provides global naming consistency for shared storage?

A.the CCR

B.the RGM

C.DID devices

D.PxFS

Answer: C

53. The cluster framework contains a kernel heartbeat module. Which two functions are provided by kernel heartbeats? (Choose two.)

A.split-brain prevention

B.detection of node faults

C.monitoring storage connectivity

D.monitoring cluster transport interfaces

E.monitoring cluster public network connectivity

Answer: BD

54. Sun Cluster 3.2 can take advantage of the new ZFS file system available with Solaris 10.

Which two new features are available with ZFS and Sun Cluster 3.2? (Choose two.)

A.ZFS has a built-in volume manager.

B.ZFS can be used as a server replication method.

C.The disks are used as a storage pool.

D.ZFS can be used as a highly available global file system.

Answer: AC

55. Which are the official locations to look up the new Solaris 10 features for Sun Cluster 3.2? (Choose two.)

A.http://docs.sun.com

B.Sun Cluster man pages

C.http://www.google.com

D.http://sunsolve.sun.com

Answer: AB

56. Which statement describes SMF as a new Solaris 10 feature?

A.SMF has completely replaced /etc/rc?.d boot scripts.

B.SMF allows the use of zones within a Sun Cluster environment.

C.SMF is required in order to boot Solaris 10 on X64 systems.

D.SMF specifies dependencies and ordering between services and provides a framework for restarting services in case of failure.

Answer: D

57. Which is an advantage of Sun Cluster 3.2 integration with SMF?

A.Sun Cluster will always start faster.

B.Sun Cluster uses SMF to shut down and reboot nodes.

C.The administrator can start/stop Sun Cluster using the svcadm command.

D.An application managed by SMF need not be rewritten to be used as a cluster service.

Answer: D

58. Which configuration can be used with ZFS and Sun Cluster 3.2?

A.ZFS as a global file system.

B.ZFS as a failover file system.

C.ZFS as a multi-owner file system.

D.A ZFS file system in /etc/vfstab (setting the file system mount point to 'legacy' within ZFS).

Answer: B

59. Solaris 10 uses SMF as a standard method of managing services.

Which two are applicable? (Choose two.)

A.Each SMF service has its own log file for diagnosing startup problems.

B.An application can be made to fail over by configuring an SMF probe within Sun Cluster.

C.SMF gives the administrator the ability to restart Sun Cluster without having to reboot the server.

D.SMF-specific proxy agents can be used to turn SMF-based applications into clustered applications.

Answer: AD

60. You have 12 ZFS file systems within the same ZFS zpool. You need them all to fail over together from node to node along with your application.

How should you configure your HAStoragePlus resource or resources to control this storage?

A.Configure 12 separate HAStoragePlus resources, one for each mount point.

B.Configure a single HAStoragePlus resource, with the Zpools property set to the name of the pool.

C.Configure 12 separate HAStoragePlus resources, each with the Zpools property set to the name of the same pool.

D.Configure a single HAStoragePlus file system with the FilesystemMountPoints property set to the list of 12 mount points.

Answer: B

61. You have a two-node cluster with a failover UFS file system that you control by setting up a HAStoragePlus resource with the FilesystemMountPoints property set to the value /oracle:/shmoracle.

Where will the file system be mounted?

A.On /oracle on nodeA but on /shmoracle on nodeB.

B.On a single mount point literally named /oracle:/shmoracle.

C.On /oracle, but if that mount point is not available the mount will fail over to /shmoracle.

D.On /oracle on either node's global zone, but loopback mount made to /shmoracle within any non-global zone on which the resource group is configured.

Answer: D

62. You want to create separate non-global zones on different nodes, and be able to fail over an application between the two zones.

What would be the advantage of using the same zone name on different nodes?

A.You can set up failover between non-global zones on different nodes only if those zones have the same name

B.You can use the syntax -n node1,node2 -z zonename only if the zone has the same name on both nodes.

C.The zones can get simultaneous access to data that is in a global file system only if they have the same name.

D.If the zones have the same name, then you will be able to switch the resource group to the zone on a specific node by simply typing

clrg switch -n nodename rgname

Answer: B

63. Cluster framework daemons are started with boot scripts on Solaris 9 and with SMF on Solaris 10.

Which statement is true?

A.On Solaris 10, you need to manually run svcadm enable for the cluster framework daemons after the cluster is initially configured.

B.On Solaris 10, every time a node boots it will NOT join the cluster automatically. You type svcadm enable for each daemon to have it join the cluster.

C.On Solaris 9, you can run /etc/init.d/clusterdaemons stop to manually disable the cluster framework on a running node. On Solaris 10 there is no such operation.

D.The daemons start automatically at cluster boot time regardless of whether you are running Solaris 9 or Solaris 10.

Answer: D

64. As cluster administrator, you need to run an application that fails over between non-global zones, and you will have a LogicalHostname resource to control IP address failover between the zones. The application needs to be able to resolve the IP address in order to run properly.

Where do you need to include the /etc/hosts entry for the IP address, assuming you are using no other name service?

A.In the non-global zones only.

B.In the global zones only.

C.In both the global zones and non-global zones.

D.In the global file system.

Answer: C

65. Sun Cluster 3.2 software provides support for applications that run in non-global zones.

Which two statements are correct? (Choose two.)

A.Sun Cluster 3.2 manages non-global zones through SMF.

B.Non-global zones are used as virtual nodes in a resource group's nodelist.

C.The Sun Cluster HA for Solaris Containers data service can fail over non-global zones between cluster nodes.

D.Sun Cluster 3.2 cannot directly manage the non-global zone or the application running in the non-global zone. This is purely a coexistence feature.

Answer: BC

66. Under what circumstances must the cluster administrator have console access to a cluster node?

A.to install cluster software

B.to monitor cluster operation

C.to repair a node that is failing to boot

D.to perform a switchover of a cluster resource group

Answer: C

67. In order to use the Cluster Console software on a management workstation, which two files must be created? (Choose two.)

A./etc/system

B./etc/clusters

C./etc/inet/hosts

D./etc/serialports

E./etc/nsswitch.conf

Answer: BD

68. Using the cluster console software, which two protocols can be used to connect to a cluster node? (Choose two.)

A.ftp

B.http

C.rlogin

D.telnet

Answer: CD

69. What is the name of the cluster console package that the cluster administrator will load onto the administrative workstation?

A.SUNWsc

B.SUNWccon

C.SUNWsccon

D.SUNWconsole

Answer: B

70. Before installing Sun Cluster 3.2 software with Solaris Volume Manager on the nodes that will make up the cluster, which three requirements must be met? (Choose three.)

A.All root partitions must be the same size.

B.The /usr file system must be on a separate partition.

C.A /globaldevices partition must be created.

D.A minimum of 750 Mbytes of swap must be configured.

E.Solaris Volume Manager requires a separate partition for its meta-state databases.

Answer: CDE

71. Which two statements identify the correct boot device restrictions? (Choose two.)

A.There must be a minimum 512 Mbytes /globaldevices file system.

B.There must be a minimum 512 Gbytes /globaldevices file system.

C.You can use a boot device that is located in a SAN, if it is only accessible from that node.

D.You can use a shared storage device as a boot device. If a storage device is connected to more than one host, it is shared.

Answer: AC

72. The cluster administrator wants to disable the BREAK signal on the serial port.

Which file needs to be edited?

A./etc/hosts

B./etc/nobreak

C./etc/default/kbd

D./etc/setkeyswitch

Answer: C

73. What is the one signal the cluster administrator never wants a terminal server connected to the serial ports of the cluster nodes to spuriously issue?

A.exit

B.login

C.break

D.reboot

Answer: C

74. A service engineer is at a customer site to install a two-node Sun Cluster 3.2 cluster on SPARC systems, and discovers that the customer has no shared storage available.

What should the engineer do?

A.Ask the customer to buy storage, since this is NOT a supported configuration.

B.Try to set up a cluster using an X64 Solaris node, to be used as third node just to have enough votes.

C.Ask the customer to provide a Solaris server (SPARC or X64) to be used as quorum server and verify with the customer that the application does not require shared data.

D.Use the Terminal Concentrator as a quorum device and install the cluster, because the customer application is purely compute-based and does NOT need shared data for the application.

Answer: C

75. The quorum mechanism and voting system are important parts of Sun Cluster.What is the majority formula to have a working cluster?

A.All the votes must be present.

B.At least 2 votes are required to form a cluster.

C.There must be a majority (more than 50 percent of all possible votes) to form a cluster or remain in a cluster.

D.Sun Cluster 3.2 automatically recalculates votes depending on the current configuration, so there is no need to worry about it.

Answer: C

76. A customer called saying their cluster is NOT starting. The customer explained that they halted node two for maintenance, then made some modification to the cluster configuration on node one. Then node one crashed. They tried to boot node two but the node is NOT allowed to form a cluster.

What is the result?

A.Split brain occurred. Check private network connection.

B.Cluster amnesia was being prevented. The customer needs to boot node one before node two.

C.Failure fencing is preventing node two from booting in order to preserve data integrity.

D.Disconnect the private network in order to disable node one, which is preventing node two from booting.

Answer: B

77. When a node leaves the cluster as a result of a failure, it is prevented from having further access to shared

storage.

What is this process called?

A.failure fencing

B.data replication

C.cluster amnesia

D.SCSI-1 reservation

Answer: A

78. Sun Cluster 3.2 introduces a new kind of quorum device called Quorum Server.

What is the minimum needed in order to use it?

A.A Linux system with Sun Cluster 3.2 packages installed.

B.A dedicated Solaris SPARC or X64 server with an scqsd daemon for each cluster.

C.A single Solaris SPARC or X64 server with one scqsd daemon for all the available clusters.

D.A Solaris SPARC server with an sql daemon and Sun Cluster 3.2 nodes with an sql daemon.

Answer: C

79. If a cluster which uses a quorum server is operating normally, and the system where the quorum server process is running is rebooted, what steps must a system administrator take?

A.Remove and re-add the quorum device.

B.None, the cluster will automatically reconnect.

C.Add the quorum device again using the clq add command.

D.Force a cluster reconfiguration event in order to connect to the new process.

Answer: B

80. You have a customer with a two-node cluster. The only storage accessible from both nodes is a NetApp NAS device. What kind of quorum configuration should the customer choose?

A.They must use a NetAPP NAS quorum device only.

B.They must use a quorum server quorum device only.

C.They can choose either the NetAPP NAS quorum device or the quorum server quorum device. If they choose both, they have too many quorum votes.

D.They should configure both a NetAPP NAS quroum device and a quorum server quorum device. This will give

them the correct number of total quorum votes and the best redundancy.

Answer: C

81. In a two-node cluster, you specifically want to assign the shared disk with the highest c#t#d# as the quorum device.

Which two methods will allow you to get this disk properly assigned as the quorum device? (Choose two.)

A.Specify which disk you want assigned as the quorum device as you run scinstall.

B.Choose during scinstall to disable automatic quorum device selection. Once both nodes are booted into the cluster, use clsetup to choose the quorum device and to reset installmode.

C.Allow scinstall to automatically configure the quorum device. Once both nodes are booted into the cluster, use clsetup first to add a second quorum device and then to remove the original one.

D.Allow scinstall to automatically configure the quorum device. Once both nodes are booted into the cluster, use clsetup first to remove the original quorum device and then to add your chosen one.

Answer: BC

82. You have a three-node cluster in a "Pair +1" configuration. Two nodes are connected to shared storage and the third is not. You choose only one quorum disk device.

What is the problem?

A.A Pair +1 cluster requires both a quorum disk device and a quorum server quorum device

B.There is no problem. Since two nodes are connected to the shared storage device, you should use exactly one quorum disk device.

C.The problem is in the "Pair + 1" configuration. Since your third node has no storage connected it is incapable of running clustered applications anyway.

D.Even with one quorum disk, you will still only be able to run the cluster if two nodes are present. But this would be the same if you had no quorum devices at all.

Answer: D

83. Which two are required for the local disks of your cluster nodes before you do an initial installation of Sun Cluster 3.2? (Choose two.)

A.You must have the nologging option for the root file system.

B.You must have swap space on each node that is twice the size of your RAM.

C.You must have the same OS revision and update installed on all nodes that are to become members of the same cluster.

D.You must have a free partition of at least 512 Mbytes on each node which scinstall will use for the /global/.devices/node@# file system.

E.If you have mutipathed storage and intend to use Veritas Volume Manager, your storage devices must already be under control of Veritas DMP.

Answer: CD

84. When setting up a three node, shared all topology using Sun Cluster 3.2, the cluster administrator wants to use a shared disk as the quorum device.

How many quorum disk devices should be installed and how many votes will each quorum device provide?

A.One quorum disk device providing two votes.

B.One quorum disk device providing three votes.

C.Two quorum disk devices providing one vote each.

D.Three total quorum disk devices, one vote from each disk device.

Answer: A

85. What path should a system administrator add to the MANPATH variable to find the manual pages for the commands in the Sun Cluster 3.2 CLI?

A./usr/share/man

B./opt/cluster/man

C./usr/cluster/man

D./usr/cluster32/man

Answer: C

86. Which three pieces of information are required to perform an interactive Sun Cluster configuration using scinstall while choosing the automatic quorum configuration? (Choose three.)

A.name of the cluster

B.name of all the nodes

C.name of all logical hostnames

D.DID number for the quorum device

E.name of the transport adapters on at least the first node

Answer: ABE

87. Which two steps are required to be able to boot your servers into the cluster? (Choose two.)

A.mirror your boot device

B.configure IPMP groups

C.select a quorum device

D.install the framework packages

E.configure the Sun Cluster 3.2 software with scinstall

Answer: DE

88. After you install a two-node cluster and reboot, before the cluster administrator assigns a quorum device, what state is the cluster in?

A.bootmode

B.installmode

C.clustermode

D.startupmode

Answer: B

89. The cluster administrator has installed a two-node cluster, reset the installmode, and chosen NOT to specify a quorum device.

If the administrator does NOT specify a quorum device, what happens to the cluster the next time one of the nodes is rebooted?

A.nothing

B.both nodes panic

C.the other node panics

D.the cluster prompts you to assign a quorum device

E.the cluster automatically assigns a quorum device for you

Answer: C

90. Which are three reasons to disable automatic quorum device configuration in a two-node cluster? (Choose

three.)

A.You have a dual ported disk or LUN that is not capable of being a quorum device.

B.You want to use a second internal drive for your quorum device.

C.You want to use the quorum server as the quorum device.

D.You want to use a NAS device as a quorum device.

E.You want to use your boot device as your quorum device.

Answer: ACD

91. During a custom scinstall when using switches for your private transport you need to assign them names. What names must you use?

A.It does NOT matter what the names are.

B.The switches do NOT require names.

C.Sun gives you the switch name.

D.The switch names must match the interface names.

Answer: A

92. Before you can configure your cluster, what must you do on each node to prepare for the global devices device files?

A.You do not do anything.

B.You must have unused space on the SAN.

C.Global devices come from your root partition.

D.You must have an unused partition on your boot device.

Answer: D

93. Which two are prerequisites for installing Sun Cluster 3.2 software with the Java Enterprise System installer? (Choose two.)

A.All Java ES auxilliary software must be installed.

B.Solaris software must be installed and correctly patched.

C.The system boot disks must be correctly partitioned for the volume manager to be used.

D.A graphical interface display must be available on the node where installation is being done.

Answer: BC

94. Which two name service lookup entries are modified with an entry of cluster in /etc/nsswitch.conf for Sun Cluster 3.2? (Choose two.)

A.hosts

B.group

C.passwd

D.ethers

E.netmasks

Answer: AE

95. A system administrator has a two-node cluster to which it becomes necessary to add a third node. Once the hardware has been connected, which two steps are required to add the new node to the cluster? (Choose two.)

A.Shut down and reconfigure all cluster nodes.

B.Run scinstall on the third node to add it to the existing cluster.

C.Use the cluster set command to modify the cluster properties.

D.On one existing cluster node use the clsetup or claccess command to add the new hostname to the authorized node list.

Answer: BD

96. Sun Cluster 3.2 has several required auxiliary components. Some of them need to be upgraded before you use the Java ES installer to install the Sun Cluster software.

How do these get installed or upgraded correctly?

A.All components are automatically installed or upgraded when you use the Java ES installer.

B.The required auxiliary components need to get installed with the clsetup command prior to starting the Java ES installer.

C.The required auxiliary components need to get installed with the scinstall command prior to starting the Java ES installer.

D.The required auxiliary components need to get manually installed with the pkgadd command. If an older package version is already installed, it needs to get manually removed with the pkgrm command before performing the pkgadd of the new version. Then you can proceed with starting the Java ES installer.

Answer: A

97. A system administrator has received the Java ES distribution as a two-CD set, and wishes to configure an NFS server to use as an install server for Sun Cluster 3.2.

What is the correct procedure?

A.Mount and share each CD from a separate server.

B.This is NOT possible, only the DVD distribution can be used in this case.

C.Create a working directory, share it using NFS, and combine the contents of both CDs into it using the cpio command.

D.Mount the first CD and share it using NFS. When prompted by the installer, unshare and unmount the CD, and repeat the process with the second CD.

Answer: C

98. If you administer a cluster with HA-Oracle you should understand some principles of the HA-Oracle fault monitor.

Which part is NOT performed by the Oracle fault monitor?

A.The HA-Oracle agent does some table manipulation.

B.The HA-Oracle agent checks some Oracle views for database activity.

C.The HA-Oracle agent creates the fault monitor user automatically, if the user is not present.

D.The HA-Oracle agent scans the Oracle Alert log for specific Oracle errors.

Answer: C

99. The cluster administrator can use scinstall interactively to configure Sun Cluster 3.2.

What are two valid options? (Choose two.)

A.configure with no questions

B.configure using GUI interface

C.configure entire cluster at once

D.configure multiple clusters at once

E.configure cluster nodes one at a time

Answer: CE

100. With Sun Cluster 3.2 it is possible to specify typical or custom installation.

Which two items must you provide to scinstall when perfoming a typical configuration? (Choose two.)

A.cluster name

B./globaldevices size

C.cluster transport adapters

D.quorum disk name

E.cluster interconnect network address and netmask

F.number of cluster interconnects to be used and switch configuration

Answer: AC

101. After the cluster administrator installs the cluster software, it has to be configured.

Which method is invalid?

A.Configure all nodes at one time.

B.Configure the nodes one by one.

C.Build various groups of nodes and configure each group one at a time.

D.Configure a single node cluster.

Answer: C

102. During the scinstall process there are many steps executed automatically.

Which two execute automatically when using scinstall on a Solaris 9 system? (Choose two.)

A.creation of the /etc/notrouter file

B.configuration of the resource groups

C.configuration of the volume management

D.selection of a quorum device in a three-node cluster

E.modification of the nsswitch.conf file with the necessary values

Answer: AE

103. After the cluster administrator has installed the Sun Cluster 3.2 packages, you configure the Sun Cluster 3.2 framework. This can be done in various ways.

Which configuration method does NOT work?

A.Using clsetup to configure the framework.

B.Use Jumpstart to both, install the packages and configure the cluster.

C.Using scinstall to configure one node at a time.

D.Using scinstall to configure the whole cluster a once.

Answer: A

104. The scinstall command can perform automatic quorum configuration. In which case must you manually configure a quorum device if you want one?

A.any cluster

B.two-node cluster

C.single-node cluster

D.any cluster with more than two nodes

Answer: D

105. When installing Sun Cluster 3.2 using the Java ES installer, two configuration options are offered:

Configure Now

Configure Later

In what circumstances is it appropriate to select the first option, Configure Now?

A.If Solaris 10 is already fully installed and patched.

B.If the system being installed is a single-node cluster.

C.When installing a completely new cluster for the first time.

D.Never, Sun Cluster 3.2 must always be configured separately after installation of the packages.

Answer: D

106. Sun Cluster 3.2 provides two management mechanisms, a Command Line Interface (CLI) and the Sun Cluster Manager, a web-based Graphical User Interface (GUI).

Which two statements about Sun Cluster Manager are true? (Choose two).

A.Sun Cluster Manager provides all the same functionality as the CLI.

B.Sun Cluster Manager provides drag and drop configuration functionality.

C.Sun Cluster Manager requires that the user be logged-in as the superuser ('root').

D.Sun Cluster Manager can be used for management, but NOT for initial cluster configuration.

E.Sun Cluster Manager can be used to continuously monitor cluster operation and performance.

Answer: DE

107. A system administrator wants to use the Sun Cluster Manager graphical interface to manage a cluster node whose hostname is node1.company.com.

What URL must be used to connect to the node?

A.http://node1.company.com/

B.https://node1.company.com/

C.http://node1.company.com:6789/

D.https://node1.company.com:8080/

E.https://node1.company.com:6789/

Answer: E

108. A system administrator is preparing to install a new cluster.

What must be consulted to determine the most recent information about required patches?

A. http://www.sun.com/bigadmin/

B. http://sunsolve.sun.com/

C.The Sun Cluster 3.2 Administration guide on the distribution medium

D.http://www.sun.com/cluster/

Answer: B

109. When you choose a transport adapter while running scinstall, you are asked:

Why would you answer no?

A.Because you want to use the same physical adapter to hold the traffic for two VLAN's, both on the private net.

B.Because your adapter will be used only for the cluster transport.

C.Because your adapter will use the tagged VLAN feature to support both private and public network traffic.

D.Because your adapter will not be dedicated to the cluster heartbeat, but will also support global file system traffic between nodes.

Answer: C

110. What are the requirements for synchronizing the time-of-day clocks between cluster nodes?

A.You must use the supplied /etc/inet/ntp.conf.cluster file so that the nodes synchronize only between themselves. You cannot use any other configuration.

B.You may use any NTP configuration, as long as the clocks on the nodes remain synchronized. The cluster will

use /etc/inet/ntp.conf.cluster file if you have no other configuration.

C.You must have both a traditional /etc/inet/ntp.conf file for all nodes to synchronize with a trusted source on the public network, and an /etc/inet/ntp.conf.cluster so that the nodes synchronize among themselves.

D.You must manually modify the /etc/inet/ntp.conf.cluster file and choose a node that will be marked with the server keyword. You must manually mark all other nodes with the client keyword.

Answer: B

111. One of the cluster nodes is at the ok boot prom prompt. You want to boot that node to multiuser and into non-cluster mode.

Which command will accomplish this?

A.init 6

B.boot -x

C.boot -sx

D./etc/reboot -- -x

Answer: B

112. Sun Cluster 3.2 uses an object-oriented command line interface (CLI) to manage the cluster.

The cluster administrator is on-site with the customer and wants to see how the cluster was configured.

Which command will be used to display the cluster configuration?

A.clnode list

B.cldevice list

C.cluster show

D.cluster status

Answer: C

113. If the cluster administrator wants to restart a process, they have to be sure it is NOT one of the cluster daemons.

Which two daemons listed below can you send a kill -HUP signal to without affecting Sun Cluster 3.2? (Choose two.)

A.rgmd

B.rpc.pmfd

C.inetd

D.syslogd

E.rpc.fed

Answer: CD

114. Which command is used to offline a resource group named nfsrg?

A.clrg switch -n node1 nfsrg

B.cluster offline nfsrg

C.clrg offline nfsrg

D.clrsdg disable nfsrg

Answer: C

115. Which command will shut down all nodes of the cluster?

A.init 0

B.shutdown -y -g0

C.cluster shutdown -y -g0

D.scadmin shutdown -y -g0

Answer: C

116. There are various ways to interact with Sun Cluster 3.2 to do daily administration.

Which method is prohibited?

A.using the clsetup command

B.editing the ccr

C.using Sun Cluster Manager

D.using standard cluster commands

Answer: B

117. Which two commands are recommended ways to shut down a node? (Choose two.)

A.halt

B.init 0

C.shutdown -y -i0 -g0

D.cluster shutdown -y -g0

E.scrgadm shutdown -y -g0

Answer: BC

118. You are talking to a customer on the phone and they need support. One of the first things you want to know is the version of Sun Cluster that they are running.

Which command do they use in order to determine the cluster version?

A.cluster show

B.cluster status

C.cat /etc/release

D.cat /etc/cluster/release

Answer: D

119. For how many clusters can one quorum server be used as a quorum device?

A.For only one cluster.

B.For as many clusters as you like.

C.From two to 16 clusters maximum.

D.None, you cannot use a server as a quorum device.

Answer: B

120. A customer has a two-node cluster with one quorum server. They have shut down node1 for maintenance. Later in the day they bring down node2 and the quorum server.

The customer has now decided to bring the cluster back up. Which answer best describes the proper technique?

A.Clear the quorum server keys, then boot node2.

B.Clear the quorum server keys, then boot node1.

C.Boot the quorum server first and then node2.

D.Boot the quorum server first and then node1.

Answer: C

121. Occasionally you may need to rebuild your cluster. If you rebuilt a cluster named "dev" without removing the old quorum server information, you may want to clean up the old quorum server information manually.

Which command would you use to remove the old quorum server information after using the command clqs show to obtain the cluster name and ID?

(--- Cluster dev (id 0x448DE82F) Registrations ---)

A.clclust remove -F quorum-serv -v

B.clqs clear -c dev -I 0x448DE82F 9000

C.scquorum -c dev -I 0x448DE82F 9000

D.Reboot the quorum server

Answer: B

122. As a precautionary measure the cluster administrator wants to have a backup of the cluster configuration. Which command can be used to provide a usable backup of the entire cluster configuration?

A./usr/cluster/bin/cluster status > /var/tmp/configfile

B./usr/cluster/bin/clrg export -o /var/tmp/configfile.xml

C./usr/cluster/bin/cluster export -o /var/tmp/configfile.xml

D./usr/cluster/bin/cluster create -i /var/tmp/configfile.xml

Answer: C

123. A patch has been released for a recently discovered security issue. The cluster administrator has downloaded the patch and the README shows that it is a rebooting node patch that must be installed in single-user mode. What steps should be followed to apply this patch with the minimum amount of downtime for the service.

A.On each node in turn run the commands:

clnode evacuate -n

shutdown -g0 -y -i0

boot -sx

Apply the patch.

Reboot back into cluster.

B.On one node run the following commands:

clrg status

clrg offline +

On each node in turn run the commands:

shutdown -g0 -y -i0

boot -sx

Apply the patch.

Reboot back into cluster.

When all nodes are patched run the following commands from one node:

clrg online +

clrg status

C.Connect to all nodes using the cluster console and apply the patch simultaneously.

Reboot each node in turn.

D.On one node run the commands:

clrg status

cluster shutdown -g0 -y

Reboot each node outside cluster.

boot -sx

Apply the patch on each node and reboot back into cluster.

boot

Answer: A

124. An Oracle DBA wishes to shut down a database, which is under Sun Cluster 3.2 control to perform maintenance.

Which two procedures could the cluster administrator use to enable the DBA to perform this task? (Choose two)

A.clrs unmonitor oracle-server-res

(DBA shuts down the database and performs maintenance)

clrs monitor oracle-server-res

B.clrg suspend oracle-rg

(DBA shuts down the database and performs maintenance)

clrg resume oracle-rg

C.clrg switch -n node2 oracle-rg

(DBA shuts down the database and performs maintenance)

clrg switch -n node1 oracle-rg

D.sqlplus / as sysdba

shutdown immediate

(DBA performs maintenance)

sqlplus / as sysdba

startup

E.clrs disable oracle-lsnr-res

(DBA shuts down the listener and performs maintenance)

clrs disable oracle-lsnr-res

Answer: AB

125. An application has failed to stop cleanly and the resource associated with it has been left in a STOP_FAILED state on node1.

What should the cluster administrator do after ensuring that the application is down?

A.clrs clear resourcename

B.clrs disable resourcename

clrs enable resourcename

C.clrs clear -n node1 -f STOP_FAILED resourcename

D.clrs export -o /var/tmp/resourceconfig.xml resourcename

clrs delete resourcename

clrs create -i /var/tmp/resourceconfig.xml resourcename

Answer: C

126. A three-node cluster contains three resource groups which have these Nodelists:

one-rg has a Nodelist: node1 node2 node3

two-rg has a Nodelist: node2 node3 node1

three-rg has a Nodelist: node3 node1 node2

The Failback property on each of the resource groups is set to FALSE.

Following a maintenance outage in which node1 and node2 have been shut down and rebooted into the cluster with node1 joining the cluster before node2:

Where will the resource groups be located, and which command should be run to move the resource groups to their preferred nodes?

A.one-rg will be on node1

two-rg will be on node2

three-rg will be on node3

No command will need to be run.

B.one-rg will be on node3

two-rg will be on node3

three-rg will be on node3

clrg remaster +

C.one-rg will be on node1

two-rg will be on node1

three-rg will be on node3

clrg remaster +

D.one-rg will be on node3

two-rg will be on node3

three-rg will be on node3

clrg evacuate -n node3 +

Answer: B

127. Which two are valid reasons to disable automatic quorum configuration? (Choose two.)

A.To configure a two-node cluster.

B.To use a NAS device as a quorum device.

C.To automatically reset the installmode flag.

D.To use the quorum server as a quorum device.

E.To have several dual-ported disks or LUNs configured.

Answer: BD

128. You are notified that your quorum disk device (d6) has failed in your two-node cluster. You need to set up another quorum disk device. You want to use d8 and verify that it is online.

Which sequence of commands is correct?

A.clq add d8

clq status

B.clq remove d6

clq add d8

clq status

C.clq add d8

clq remove d6

clq status

D.clq replace d6 d8

clq status

Answer: C

129. You have a three-node cluster. The node names are jupiter, saturn and mars. You need to disable disk path monitoring for d11 from all nodes.

Which two commands will perform this task? (Choose two.)

A.cldev unmonitor d11

B.cldev unmonitor +

C.cldev unmonitor -n jupiter,saturn,mars d11

D.cldevice unmonitor -n jupiter,saturn,mars +

Answer: AC

130. You have a two-node cluster that is using scsi-2 disk reservations for your shared storage. All your shared disks support scsi-3 and their default fencing property has not been changed. The cluster's global_fencing property is set to pathcount. The quorum device is d7 and you want d7 to remain as your quorum device. Your shared devices are d3-d10.

Which set of commands will change all disks to use scsi-3?

A.cluster set -p global_fencing=scsi3

clq refresh d7

B.cluster set -p global_fencing=scsi3

clq add d6

clq remove d7

C.cluster set -p global_fencing=prefer3

clq remove d7

clq add d7

D.cluster set -p global_fencing=prefer3

clq add d6

clq remove d7

cldev set -p default_fencing=global d7

clq add d7

clq remove d6

Answer: D

131. You have been asked to change the default 10-second timeout setup of the cluster heartbeat to four seconds. Which identifies the correct procedure to accomplish this task?

A.cluster set -p heartbeat_timeout=4000

B.cluster change heartbeat_timeout=4000

C.1) cluster set -p heartbeat_quantum=800

D.1) cluster set -p heartbeat_timeout=400

Answer: C

132. Which two are the default settings for Disk Path Monitoring? (Choose two.)

A.all disk paths are monitored

B.only disk paths to shared devices are monitored

C.reboot_on_path_failure is enabled

D.reboot_on_path_failure is disabled

Answer: AD

133. Identify the required steps in order to enable automatic reboot on a two-node cluster (nodes mars and earth), if all the monitored disk paths from that node are broken and there is another node that still has working paths to at least one of the shared devices.

A.1) clnode set -p reboot_on_path_failure=enabled mars

2) clnode set -p reboot_on_path_failure=enabled earth

B.1) Identify all DID numbers that represent disks only connected to one node.

2) For each device, perform cldev unmonitor d

3) clnode set -p reboot_on_path_failure=enabled mars

4) clnode set -p reboot_on_path_failure=enabled earth

C.Nothing required, this feature is enabled by default already.

D.1) Identify all DID numbers that represent disks connected to all nodes.

2) For each device, perform cldev unmonitor d

3) clnode set -p reboot_on_path_failure=enabled mars

4) clnode set -p reboot_on_path_failure=enabled earth

Answer: B

134. A two-node cluster has been configured with a scsi-2 quorum device. The administrator has been asked to create a ZFS pool using this storage device.

Which action should the administrator perform in order to configure the ZFS pool?

A.The ZFS license key should be installed on all nodes.

B.SCSI-3 reservations should be enabled on the quorum device.

C.Disk path monitoring should be enabled on the storage device.

D.The quorum device must be removed before creating the ZFS pool and the quorum re-added afterward.

Answer: D

135. A cluster administrator wants to give the user pat rights to switch the resource groups in the cluster.

Which task should the administrator perform?

A.Assign the authorization solaris.cluster.admin to the user pat.

B.Nothing, the authorization is by default part of the Basic Solaris User profile.

C.roleadd -u 1234 -d /export/home/clusoper -A solaris.cluster.modify

usermod -R clusoper pat

D.Tell pat the URL for the Sun Java Web Console and allow the user to login with their username and password.

Answer: A

136. The cluster administrator is running SunCluster 3.2 on x86 machines running Solaris 10 update 3, and needs to boot one of the x86 machines into non-cluster mode.

To which menu item does the cluster administrator add '-x' using the GRUB editor?

A.kernel /platform/i86pc/multiboot

B.kernel /platform/i86pc/suserboot

C.Solaris 10 11/06 s10x_u3wos_08 X86

D.module /platform/i86pc/boot_archive

Answer: A

137. A customer has a three-node cluster configured as a pair-plus-one topology. The customer asks about whether they need to use any quorum devices, and if so how many.

Which answer best describes how the administrator will respond?

A.You do NOT need any quorum devices.

B.You must have a quorum device, all clusters require quorum devices.

C.You do NOT need any quorum devices, although it would make your cluster more highly available. You should configure ONE quorum device for highest availability.

D.You do NOT need any quorum devices, although it would make your cluster more highly available. You should configure TWO quorum devices for highest availability.

Answer: D

138. How many votes are assigned to a quorum device?

A.All quorum devices are assigned one vote.

B.Quorum device votes are equal to the number of node votes.

C.The amount of votes assigned to a quorum device is equal to the amount of node votes connected to it.

D.The amount of votes assigned to a quorum device is equal to one fewer than the number of node votes connected to it.

Answer: D

139. The cluster has been configured using qfe0 and qfe4 as the private network adapters on each node in a two-node cluster. On node1 qfe0 is starting to generate errors and the customer would like to reconfigure node1 to use qfe1 as a replacement for qfe0. The customer does not accept downtime.

Which sequence of commands should be used to accomplish this?

A.clintr remove node1:qfe0

clintr add node1:qfe1

B.the cluster must be shut down

C.clintr disable -n node1:qfe0,node2:qfe0

clintr remove -n node1:qfe0

clintr add -n node1:qfe1

clintr add -n node1:qfe1,node2:qfe0

D.unplug the cable from node1 qfe0 and plug the cable into qfe1

Answer: C

140. A customer has a four-node cluster configured as a scalable topology. No quorum devices have been configured.

How many nodes must be present in order for the cluster to be able to offer any of its application services?

A.one

B.two

C.three

D.four

Answer: C

141. The cluster administrator has a three-node cluster and they need to perform maintenance on one of the nodes. The administrator expects the node to be out of the cluster for an extended period of time.They can disable this node's quorum vote during the maintenance period.

What is this method of operation called?

A.Unconfiguring the node from the cluster.

B.Deleting the node from the cluster.

C.Putting the node into maintenance state.

D.You cannot disable a quorum vote without crashing the cluster.

Answer: C

142. You have a three-node cluster and have put one node into maintenance mode.

What do you need to do to give the node its vote back and bring it back into the cluster?

A.Re-install the node to fix the maintenance mode.

B.Reboot the node, the maintenance mode will reset automatically.

C.Remove the node from the cluster, then add it back into the cluster.

D.Add another quorum disk device before booting the node back into the cluster.

Answer: B

143. You have a two-node cluster configured. Node one is out of the cluster. You make changes to node two (add resource groups and disk device groups). You now shut down node two and try to boot node one, which fails due to amnesia prevention.

Which two keep node one from starting the cluster? (Choose two.)

A.scsi-2 reservations

B.Persistent reservations on the quorum device.

C.Nothing will keep node one from starting the cluster.

D.Node one's key has been previously preempted from the quorum device.

Answer: BD

144. You have a scsi-2 quorum device in a two-node cluster. What happens to that quorum device if you change the global fencing policy using the command cluster set -p global_fencing=prefer3?

A.The quorum device is removed from the configuration.

B.The quorum device automatically changes to a scsi-3 device.

C.The per-disk fencing policy for the quorum disk is changed to pathcount.

D.The request to change the global fencing policy is denied, because you have a scsi-2 quorum device.

Answer: C

145. You have been running your four-node cluster with three private networks for a while and now you want to change the IP address space used for the entire cluster transport from 172.16.x.x to 192.168.5.x.

How can you accomplish this change?

A.Run clsetup on any one node in cluster mode

B.Boot all nodes to non-cluster mode and run clsetup on all nodes to change the private network information.

C.You cannot accomplish this change. It is not supported to change the IP address range used for the cluster transport.

D.Boot all nodes to non-cluster mode, enable root ssh or rsh between them, and run clsetup from one node only to change the private network information.

Answer: D

146. You want to give a non-root user the ability to online, enable, disable, and switch over elements of the cluster, but NOT to modify the configuration.

In which two ways can you accomplish this? (Choose two.)

A.Give that specific user the solaris.cluster.admin authorization.

B.Give that user the ability to su to a non-root role that has the solaris.cluster.admin authorization.

C.Give that user the ability to su to a non-root role that has the solaris.cluster.modify authorization.

D.You do not have to do anything. Every user can accomplish what you want since they will naturally all inherit the solaris.cluster.read authorization from the Basic Solaris User profile.

Answer: AB

147. Which two statements describe reasons for using the Sun Cluster Manager graphical application to manage the cluster (Choose two.)?

A.It shows live status of the cluster by highlighting faulted components, without the user having to manually refresh the display.

B.It can draw pictures of some of the hardware and other configured elements of your cluster.

C.It has its own authentication mechanism separate from the node on which it is running. Users can log in with user names different than those on the node, and thus be authorized to accomplish tasks that they cannot accomplish on the command line.

D.It contains wizards to guide you through properly configuring IPMP on your public network.

Answer: AB

148. You have a four-node cluster, where node1 and node2 are connected to some shared storage, and node3 and node4 are connected to different shared storage.

Now node1 has died and you will not be able to reboot it for some time. Why would you then want to run clq disable node1?

A.This will prevent node1 from accidentally rebooting until you explicitly run clq enable node1 on one of the remaining nodes.

B.You may want to run the command but it is actually too late. You can only run the disable option when the node mentioned is still up in the cluster.

C.This will alter the quorum vote count in the remaining cluster so that if later you also lost node2, the remaining nodes could continue running the cluster.

D.This is simply an informational flag. The output of the clq status will specifically show a maintenance flag for node1, indicating that you recognize it is dead and will not be able to reboot for some time.

Answer: C

149. You set up your cluster using the default 172.16.x.x range for the cluster transport. Yet these same IP addresses are in use on machines in a department that you need to access across the public network.

What kind of problem will you experience?

A.Your cluster nodes will refuse to communicate across the private network.

B.Your public network adapters will be marked as FAILED.

C.Communication to the machines using 172.16.x.x on the public network fails, because your cluster nodes will want to route this traffic across the private network.

D.You will remain free of problems. It is perfectly acceptable to use a range of IP addresses on your private networks that conflicts with other machines on your public network.

Answer: C

150. You have been asked to change the network addressing and ranges for the cluster transport from the default values to use the subnet 192.168.42.0, with anticipated future growth to eight nodes and an anticipated number of four private networks.

Which identifies the correct steps to follow?

A.Run clsetup and select "Change Network Addressing and Ranges for the Cluster Transport". Answer the questions provided.

B.1) Shut down the cluster and boot all nodes in single-user non-cluster mode.

2) Run clsetup and select "Change Network Addressing and Ranges for the Cluster Transport". Answer the questions provided.

3) Reboot the nodes into multi-user cluster mode.

C.1) Shutdown the cluster and boot all nodes in multi-user non-cluster mode.

2) Run clsetup and select "Change Network Addressing and Ranges for the Cluster Transport". Answer the questions provided.

3) Reboot the nodes into multi-user cluster mode.

D.1) Shutdown the cluster and boot all nodes in single-user non-cluster mode.

2) Modify the cluster infrastructure file on all nodes and amend the necessary properties as specified.

3) Reboot the nodes into multi-user cluster mode.

Answer: C

151. You want to create a role that contains only the authorizations solaris.cluster.modify and solaris.cluster.admin. As you then decide which users should be able to successfully access every cluster command, you can just assign these users to the role and give them the role password. These users will never need to have the root password, but they will be able to execute every cluster command.

Which two statements are true? (Choose two.)

A.This scheme will work for both cluster commands and Sun Cluster Manager.

B.There are still some cluster-related activities, such as modifying the /etc/vfstab file, that the user might NOT be able to do.

C.The users who switch to the role will still NOT have solaris.cluster.read authorization, so they will not be able to view cluster status and configuration.

D.This scheme will work as the users try to access command line commands, but it will NOT work if the users execute operations through Sun Cluster Manager.

Answer: AB

152. Which commands would be used to define a third private network using switches for a two-node cluster?

A.clintr add-privnet -s switch3 node1:bge2,node2:bge2

B.clintr add node1:bge2

C.clintr add node1:bge2

D.clintr add -s switch3 node1:bge2 node2:bge2

Answer: B

153. How can you name Solaris Volume Manager disksets that include disks from shared storage?

A.You can name them anything you want, often related to what application you may be using them for.

B.The name of the diskset must be the same as the name of the cluster.

C.The Solaris Volume Manager software automatically chooses a name for the diskset, such as ds101 or ds102. You must keep this name.

D.A diskset that is physically attached to nodes named vincent and theo must be named ds_vincent_theo, and there can be only one such diskset.

Answer: A

154. You have two local (non-shared) disks on a cluster node. You put one local metadb replica on one disk and

two local metadb replicas on the other disk.

What will happen if later the disk with the two replicas fails?

A.You are guaranteed that the node stays up and running, and can reboot. As long as Solaris Volume Manager can find one valid copy of the configuration, the node will stay up.

B.You are guaranteed to stay up and running. However, if you reboot, you may have to manually delete the broken metadb replicas before being able to join the cluster.

C.When Solaris Volume Manager discovers you have less than 50% of your local metadbs remaining, it will fail immediately.

D.When Solaris Volume Manager discovers you have less than 50% of your local metadbs remaining, it will prompt you to fix the broken ones and you can stay operational without rebooting.

Answer: C

155. How do you define Solaris Volume Manager volumes, when configuring your boot device?

A.You always define the physical disk partition as a submirror, using the standard c#t#d#s#.

B.You always define soft partitions, so that you can grow your root file system.

C.You always define your submirror using the DID device /dev/did/rdsk/d#, so that your volumes are usable in the cluster.

D.You always make sure both submirrors are attached to the mirror before editing the /etc/vfstab file to mount the mirror devices.

Answer: A

156. How is a disk repartitioned when it is added to a diskset?

A.The entire drive is mapped to slice 0.

B.A small portion of the drive is mapped to slice 3 and the remainder to slice 4.

C.A small portion of the drive is mapped to slice 7 or 6 and the remainder to slice 0.

D.The drive is automatically converted to an EFI label regardless of the size of the drive.

Answer: C

157. Which is important when installing Sun Cluster with Solaris Volume Manager to manage the boot disks?

A.The /global/.devices/node@1, /global/.devices/node@2 should have unique volume names for each node in the cluster.

B.All volume names should be configured to be the same across all cluster nodes.

C.The /global/.devices/node@1, /global/.devices/node@2 should be configured on a soft partition.

D.The disk partition 7 on the boot disk should be configured to be 6 Mbytes for the metadbs.

Answer: A

158. You want to configure five disksets within a cluster environment running on Solaris 10. What must you do for Solaris Volume Manager to allow you to create more than four disksets?

A.Nothing is required, the number of disksets is dynamically updated in Solaris 10.

B.The maximum number of disksets that can be configured in Solaris 10 is four.

C.Each diskset must be registered into the cluster environment using cldg sync

D.Adjust the parameter md_nsets in /kernel/drv/md.conf to allow more than four disksets to be created.

Answer: A

159. Which two are requirements for Solaris Volume Manager shared storage disksets in Sun Cluster 3.2? (Choose two.)

A.They must contain four or more disks.

B.They can be attached to a maximum of two nodes.

C.They must contain disks only in the shared storage arrays.

D.They must contain a combination of non-shared disks and shared disks.

E.If they contain disks that are in more than one array, then those arrays must be attached to exactly the same set of nodes.

Answer: CE

160. A cluster administrator has yet to run a single Solaris Volume Manager command since they installed Solaris. Now the administrator has configured the cluster framework and intends to use Solaris Volume Manager for the shared storage.

What is the very first Solaris Volume Manager task that they must do?

A.Create diskset mediators.

B.Manually add metadb replicas on shared storage disks.

C.Create empty disksets by specifying the names of nodes connected to the shared storage.

D.Manually add local metadb replicas on local disks of each node connected to the shared storage.

Answer: D

161. You just created a Solaris Volume Manager diskset from the shared storage containing two disks, one from each of two arrays. You added two connected nodes as mediators for the diskset.

Which two of the following are true? (Choose two.)

A.You must manually create a partition on the diskset disks to hold a metadb replica.

B.The metainit commands that you use to build volumes on the disks will fail unless you mirror across arrays.

C.The mediators allow you to lose a disk and subsequently maintain high availability if a node fails.

D.If you lose one of the disks and then lose the node that is primary for the diskset, the entire cluster will lose access to the diskset.

E.The Solaris Volume Manager commands that added disks to the diskset automatically created a single metadb replica on each disk.

Answer: CE

162. The metaset commands that you used to build a diskset have already registered the diskset as a cluster device group. For which two reasons would you still use the cldevicegroup (cldg) command for this diskset? (Choose two.)

A.You would use cldg sync if you add or delete a volume.

B.You would use cldg create to restore the diskset to the CCR if it fails.

C.You would use cldg switch to manually switch the primary for the device group.

D.You would use cldg set to change properties of the device group, such as preferenced and failback

Answer: CD

163. You run cldg status to view the status of your device group. What must be true about the node that is listed as the Primary?

A.It is the preferred node for the device group.

B.It is the current primary node for the device group, not necessarily the preferred node.

C.A global file system configured on a volume of the device group can only be accessed from that node.

D.The device group can run only on that node. If that node fails, the device group will become unavailable.

Answer: B

164. A two-node Solaris Volume Manager campus cluster has been installed at sites A and B. A preferred site A can be configured so that the cluster survives the loss of site B without requiring manual intervention.

Which is the correct procedure to create a preferred site A ?

A.Put the quorum device at site A and configure the majority of the diskset metadb replicas on the disks located at site A.

B.Put the quorum device at site A and add an extra metadb to the disks in the diskset at both sites A and B.

C.Put the quorum device at site A and add an extra local metadb to the root disk on the server located at site A.

D.You cannot do this unless you configure a third site with the quorum device and with the diskset configured across all three sites, so that the loss of one site will always ensure that two out of the three metadbs are present.

Answer: A

165. A diskset has been created in a two-node cluster with the incorrect name. No disks have been added to it. How does the cluster administrator remove the diskset from the cluster framework?

A.cldg list

cldg status disksetname

cldg offline disksetname

cldg delete disksetname

cldg list

B.cldg list

cldg status disksetname

metaset -s setname -A remove

cldg list

C.cldg list

cldg status disksetname

metaset -s disksetname -d -f -h node1 node2

cldg list

D.cldg list

cldg status disksetname

cldg offline disksetname

cldg remove-device

cldg list

Answer: C

166. You have a four-node cluster with all nodes connected to the storage with a diskset disk-ds1 configured on the shared storage.

(a) What are the default number of secondaries?

(b) How do you alter the number of secondaries so that there is a primary node and the remainder are all secondaries?

A.(a) All nodes are primaries or secondaries by default.

(b) You do not need to alter anything.

B.(a) There is one secondary by default.

(b) Run the commands:

cldg set -p numsecondaries=4 disk-ds1

cldg show disk-ds1

C.(a) There is one secondary by default.

(b) Run the commands:

cldg set -p numsecondaries=3 disk-ds1

cldg show disk-ds1

D.(a) There are no secondaries by default

(b) Run the commands:

cldg set -p numsecondaries=3 disk-ds1

cldg show disk-ds1

Answer: C

167. Your customer has chosen to use Solaris Volume Manager for the shared disks. A diskset named ds1 has already been created with node1 and node2 being able to master the diskset. The customer would like to build a mirrored volume to contain their application data.

The diskset ds1 contains two 36 Gbyte disk drives using did devices d4 and d8.

The customer would like to use as volume names d101 and d102 as the submirrors and d100 as the mirror.

The customer would like to create two soft partitions named d10 and d11, with d10 5gb and d11 10gb in size.

Which sequence of commands will accomplish this task?

A.metainit -s ds1 d101 1 1 /dev/did/rdsk/d4s0

metainit -s ds1 d102 1 1 /dev/did/rdsk/d8s0

metainit -s ds1 d100 -m d101

metattach -s ds1 d100 d102

metainit -s ds1 d10 -p d100 5g

metainit -s ds1 d11 -p d100 10g

B.metainit d101 1 1 /dev/did/rdsk/d4s0

metainit d102 1 1 /dev/did/rdsk/d8s0

metainit d100 -m d101

metattach d100 d102

metainit d10 -p d100 5g

metainit d11 -p d100 10g

C.metainit d101 1 1 /dev/did/rdsk/d4s0

metainit d102 1 1 /dev/did/rdsk/d8s0

metainit d100 -m d101,d102

metainit d10 -p d100 5g, d11 -p d100 10g

D.metacreate -s ds1 d101 /dev/did/rdsk/d4s0

metacreate -s ds1 d102 /dev/did/rdsk/d8s0

metamirror -s ds1 d100 d101,d102

metaspart -s ds1 d10 d100 5g

metaspart -s ds1 d11 d100 10g

Answer: A

168. Which command will switch a device group, named ds1, from node1 to node2?

A.cldg switch -n node2 ds1

B.haswitch -d ds1 -h node2

C.node1> metaset -s ds1 -r

node2> metaset -s ds1 -t

D.clswitch ds1 from node1 to node2

Answer: A

169. Given this output from the cldg show command:

Device Groups ===

Device Group Name:ds1

Type:SVM

failback: false

Node List:node1, node2

preferenced:true

numsecondaries:1

diskgroup name:ds1

Which command would enable the failback property?

A.cldg set -p failback=true ds1

B.cldg set -p failback=enable ds1

C.metaset -s ds1 set failback=true

D.cldevicegroup set failback=true ds1

Answer: A

170. You need to take one of your Solaris Volume Manager device groups offline for some emergency repairs. The device group is called ds1 and all devices are used as file systems. You have unmounted all the file systems.

Which command would you run to offline the device group?

A.cldg offline ds1

B.cldg switch offline

C.cldg offline -s ds1

D.cldg switch offline -s ds1

Answer: A

171. Which two statements are true about a global file system? (Choose two.)

A.It is seen as mounted only on one node.

B.All nodes have simultaneous access to the data.

C.All nodes directly connected to the storage can do direct i/o.

D.All nodes do NOT have to be directly connected to the storage.

Answer: BD

172. Which two statements are true about local failover file systems?

A.Mount at boot is set to yes in the /etc/vfstab file.

B.All nodes have simultaneous access to the file system.

C.Only one node can have the file system mounted at any one time.

D.Only nodes physically attached to the storage can mount the file system.

Answer: CD

173. Which two options in /etc/vfstab are valid for the global file system? (Choose two.)

A.Mount at boot is set to yes.

B.Mount options must be set to false.

C.The mount point must have the word global.

D.The word global is one of the mount options.

Answer: AD

174. Which two statements identify the requirements for installing VxVM 4.x in Sun Cluster? (Choose two.)

A.You must have a rootdg configured on all nodes.

B.VxVM must be installed on all nodes physically connected to the shared storage.

C.All disks must get partitioned using the Cross-Platform Data Sharing (CDS) layout.

D.The vxio major number assigned in /etc/name_to_major file must be identical on all nodes.

E.All disks must be traditionally partitioned using separate Solaris OS disk slices for the private region and the public region.

Answer: BD

175. Which two statements identify the tasks a node can perform, when that node has imported a Veritas disk group? (Choose two.)

A.The node will increase the size of the private region on the disks within the disk group.

B.The node physically reads and writes data to the drives within the disk group.

C.The node waits until all mirrored data is in sync before writting data to the disk group.

D.The node can voluntarily give up ownership of the disk group by deporting the disk group.

Answer: BD

176. What identifies the only requirement for creating a Veritas volume in a shared storage disk group for the Sun Cluster 3.2 environment?

A.You must stripe all data across all available controllers.

B.You must decide between mirroring on the VxVM layer OR within the hardware raid controller.

C.You must mirror across controllers unless there is full redundancy provided by the controller itself.

D.You must make sure that the storage devices you mirror are physically located in different data center rooms.

Answer: C

177. What is the complete list of acceptable Veritas volumes to hold your Sun Cluster 3.2 application data on arrays of JBOD disks?

A.mirrored stripe, striped mirrors

B.simple mirrors, mirrored stripe, striped mirrors

C.concatenated subdisks, striped subsets, simple mirrors

D.striped mirrors, striped subdisks, concatenated subdisks, concatenated mirrors

Answer: B

178. If you choose to encapsulate the boot disk using Veritas Volume Manager 4.x and above, you must create a disk group that will contain the disk.

Which two statements are correct? (Choose two.)

A.The the name of the disk group cannot be rootdg.

B.The name of the disk group must be rootdg.

C.You can choose any name for the group that you would like except bootdg.

D.Veritas will use the reserved name bootdg to link to the name that you provided.

Answer: CD

179. Your customer has a fully redundant hardware RAID storage environment. Your customer has also chosen to place the hardware RAID LUNs under Veritas Volume Manager control.

Which statement is correct?

A.Software RAID always has better performance.

B.You need to add additional RAID protection to the LUNs

C.Hardware RAID does NOT provide total protection of your data.

D.You can put the LUNs under Veritas Volume Manager control for the purpose of space management.

Answer: D

180. Your customer has asked you to create a Veritas Volume Manager volume. They want the volume to be RAID level 1+0. The name of the group is oradg and they want the name of the volume to be oravol, and 3 Gbytes in size.

Which three statements are correct? (Choose three.)

A.RAID 1+0 requires at least 4 disks

B.The volume oravol will be a layered volume

C.vxmake -g oradg oravol layout=mirror

D.vxassist -g oradg make oravol 3g layout=stripe-mirror

E.vxassist -s oradg make oravol 3g layout=mirror-stripe

Answer: ABD

181. Your customer has just created a new Veritas Volume Manager volume named oravol in the already registered disk group oradg. Now they are trying to create a UFS filesystem on top of the volume and keep getting the following error message after executing the command:

newfs /dev/vx/rdsk/oradg/oravol

/dev/vx/rdsk/oradg/oravol: No such device or address

Which two statements describe the cause of the error? (Choose two.)

A.You did not run cldg sync oradg.

B.You typed the newfs command incorrectly.

C.The newfs command must be run on one of the other nodes.

D.The global namespace does NOT know about this new volume on the other nodes.

Answer: AD

182. Which three are objects used within Veritas Volume Manager?

(Choose three.)

A.plexes

B.volumes

C.subdisks

D.metamirrors

E.diskset

F.metaplex

Answer: ABC

183. When configuring Veritas Cluster Volume Manager for use with a two-node Oracle RAC Sun Cluster 3.2 install, which node can configure the shared disk group?

A.Both nodes can configure the group.

B.Neither node can configure the group.

C.The master node can configure the group.

D.You cannot use CVM in Sun Cluster 3.2.

Answer: C

184. In Veritas Volume Manager 4.x or 5.x, once you encapsulate your boot device it goes into a disk group. This group historically was called rootdg.

What is the name of the link to this diskgroup under 4.x and 5.x control?

A.rootdg

B.bootdg

C.There is no link to this diskgroup.

D.You must link this group to rootdg.

Answer: B

185. When would you use a Veritas Volume Manager CDS disk in the bootdg group?

A.All the time, it is the default.

B.You cannot use CDS disks in the bootdg group.

C.When you want all nodes to boot from the same device.

D.When you want to share your boot device between clusters.

Answer: B

186. Which two are correct statements for the two text-based installers of the Veritas products (installer and installvm)? (Choose two.)

A.They help to automate the encapsulation of the boot disks on all nodes.

B.They help to make sure that the major number for vxio is the same across all nodes.

C.They guide you through entering licenses and initializing the software at the end of the installation.

D.They can install the software on multiple nodes if rsh/rcp or ssh/scp is enabled between the nodes.

Answer: CD

187. How do you decide what DID name to give to your Veritas Volume Manager disks as you assign them to disk groups?

A.The cluster assigns the DID name to the VxVM disk.

B.VxVM automatically assigns the DID name to the disk.

C.You have to manually assign the DID name to a VxVM disk.

D.VxVM uses its own naming scheme, it does NOT use the DID name.

Answer: D

188. Which two Veritas Volume Manager packages need to be installed on all nodes connected to the shared storage? (Choose two.)

A.VRTSjre

B.VRTSvlic

C.VRTSvxvm

D.VRTSobgui

E.VRTSvxman

Answer: BC

189. When using Veritas Volume Manager in the cluster, why would you choose to use Dirty Region Logging (DRL) for your mirrored volumes?

A.DRLs can only be used for RAID 5 volumes.

B.You would never use DRLs, they are too large.

C.You CANNOT use DRLs within Sun Cluster 3.2.

D.You should use DRLs to minimize the resynchronization of your mirrored volumes after a system crash.

Answer: D

190. Suppose you have a Veritas Volume Manager disk group named appdg with default settings registered using:

cldg create -t vxvm -n node1,node2 appdg.

You are now asked to configure the following: appdg should be online on node1 by default, and node1 should take appdg automatically back when it joins the cluster.

Which identifies the correct action to accomplish this task?

A.No action is required, since this is already the default setup.

B.cldg set -p preferenced=true -p failback=true appdg

C.cldg set -p preferenced=node1 -p failback=true appdg

D.cldg change -p preferenced=true -p failback=true appdg

Answer: B

191. You have a two-node cluster with nodes mars and venus. Given the following output:

hostname

mars

vxdg list

NAMESTATEID

rootdgenabled1152724944.22.mars

nfsdgenabled1152725899.26.venus

Note the ID of disk group nfsdg.

Which statement is correct?

A.The ID of nfsdg looks incorrect. You should contact Sun support for further investigation.

B.Disk group nfsdg was previously imported by node venus and now is imported on node mars. The ID just reports the history.

C.Disk group nfsdg is currently imported on node mars. The ID just reflects the node where the disk group was initially created. There is no problem.

D.Disk group nfsdg seems to have a problem, since it is reported to be imported on node mars, but still carries an ID from venus. So you might have to suspect data corruption.

Answer: C

192. Your customer has just successfully executed the following commands:

vxassist -g oradg make oravol 3g layout=mirror

cldg sync oradg

newfs /dev/vx/rdsk/oradg/oravol

mkdir /oracle

Now the customer is trying to mount the file system globally and gets this error message after typing the commands:

mount -o global /dev/vx/dsk/oradg/oravol /oracle

mount: /dev/vx/dsk/oradg/oravol or /oracle, no such file or directory

The customer checks the device file and the directory on that node and determines that they do exist.

What has the customer forgotten to check?

A.The command was typed incorrectly.

B.The mount point does not exist on all nodes.

C.The file /dev/vx/dsk/oradg/oravol does not exist on all nodes.

D.The customer has forgotten to synchronize the new volume information.

Answer: B

193. You need to remove the Veritas Volume Manager volume named nfsvol in the nfsdg disk group which is mounted on /nfs.

Which set of commands are used to perform this procedure?

A.umount /nfs; cldg remove nfsdg -r volume nfsvol; scconf sync nfsdg

B.umount /nfs; vxassist -g nfsdg remove volume nfsvol; cldg sync nfsdg

C.unmount -g nfsdg all; vxassist -g nfsdg remove all; cldg sync nfsvol

D.unmount /nfs; vxassist -g nfsdg remove volume nfsvol; cldg sync nfsvol

Answer: B

194. You are using Veritas Volume Manager to manage Sun Cluster 3.2 shared storage on a multipathed Hitachi array. Which two software products can you use to manage the multipathing? (choose two)

A.Veritas DMP

B.Sun StorEdge Traffic Manager Software (built into Solaris 10)

C.Hitachi Dynamic Link Manager

D.Veritas Volume Replicator

Answer: BC

195. When using ZFS in the cluster environment, you generally do not need any other volume manager to manage your data storage. Instead, you create pools of storage to manage the file systems within.

What is the ZFS storage unit called?

A.zraid

B.zstor

C.zpool

D.zunit

Answer: C

196. You have a ZFS pool containing a single disk c2t0d0. The pool contains a file system that you want to mount on /data.

What does the entry look like in the /etc/vfstab file?

A.ZFS populates the /etc/vfstab entries for you automatically

B.using ZFS, you do NOT create /etc/vfstab entries

C./dev/dsk/zpool1/dev/rdsk/zpool1 /data zfs 3 no -

D./dev/dsk/c2t0d0s0 /dev/rdsk/c2t0d0s0 /data zfs 1 yes global

Answer: B

197. When using ZFS within the cluster, the cluster administrator should mirror the data across controller paths. Each node should have at least two separate paths to shared storage.

How does the cluster administrator ensure that the paths are identical on each node so the ZFS pool can fail over without any problems?

A.Do nothing. Even if the path names are different on each node, zpools can still fail over.

B.ZFS automatically renames disk path names to ensure that they are identical on each node.

C.Use the DID paths.

D.Manually ensure that each controller has the same number on each node.

Answer: A

198. The cluster administrator is using ZFS to manage the shared storage for the cluster. You want to create a pool called "mypool".

What is the command used to create a ZFS pool, and to verify that the pool has been created?

A.1) zcreate mypool c2t0d0s0

2) zstat

B.1) zfs create mypool/myfs1

2) zpool status

C.1) zpool create mypool mirror c2t0d0s0 c3t0d0s0

2) zpool status

D.1) zpool create mypool mirror c2t0d0s0 mirror c3t0d0s0

2) zlist

Answer: C

199. When using ZFS for your shared storage within the cluster, the mountpoints default to "/poolname/fsname". You can change the mountpoint name to whatever you want. You want to change the mountpoint for the filesystem myfs1 in the pool devpool to /oradb.

Which command do you use?

A.zset mount=/oradb devpool/myfs1

B.zpool setmount=/oradb poolname/fsname

C.zfs /oradb set mountpoint devpool/all

D.zfs set mountpoint=/oradb devpool/myfs1

Answer: D

200. When using ZFS to maintain your cluster file systems, what other software volume management application do you need to manage the device groups?

A.You do not need any other volume manager software.

B.SVM is needed to manage your disksets.

C.VxVM is needed to manage your disk groups.

D.Veritas CVM is needed to manage the shared disk groups.

Answer: A

201. When managing ZFS with an HAStoragePlus resource, what extension property is used to specify the storage?

A.ZFS

B.Zpools

C.GlobalDevicePaths

D.FilesystemMountPoints

Answer: B

202. A system administrator has configured a ZFS file system in a cluster as a failover resource. In order to make this file system available to clients outside the cluster through NFS, what additional configuration is required?

A.Add a share command for the zpool to /etc/dfs/dfstab.

B.None, ZFS file systems in a cluster cannot be shared by NFS.

C.Use the standard Sun Cluster HA-NFS agent to share the file system.

D.Use the zfs sharenfs=on command to make the file system shareable.

Answer: C

203. Which is a reason for using link-state testing rather than "ping testing" for public network adapters in IPMP groups in Solaris 10?

A.Link-state testing is always more robust.

B.IPMP will do repair detection and failback only if you use link-state testing.

C.The cluster will only be able to monitor the state of IPMP groups if you use link state testing

D.You do NOT have to consume public network IP addresses for test addresses if you use link-state testing.

Answer: D

204. What is the reason that cluster still requires a PNM daemon even though IPMP itself is built into the cluster?

A.PNM is required to support multiple IPMP groups on the same subnet. Without PNM, IPMP itself can not support this.

B.The PNM daemon is required to automatically configure your adapters into IPMP groups if you have not already done so manually.

C.IPMP can cause failover of IP addresses between adapters on the same standalone server, but not on a cluster node.

D.IPMP is not cluster-aware. The PNM cluster daemon is required so that applications can fail over between nodes if all adapters in an IPMP group on one node have failed.

Answer: D

205. You have two adapters that you want to configure in the same IPMP group. Neither of them are yet configured into any IPMP group. The cluster nodes are already configured.

Which tasks must you perform to configure the group?

A.You need to rerun the scinstall utility and let it create your IPMP groups for you.

B.You need to create or edit the /etc/hostname.xxx files for each adapter, using the group option to place them in the same group.

C.You need to use the clinterconnect (clintr) command and specify the adapters for which you want to create IPMP groups.

D.You need to create or edit the /etc/hostname.xxx file for only one adapter. In that file, you list all the adapters that will be in the same group.

Answer: B

206. When you design the network redundancy for a cluster node, you have to understand the key requirements of IPMP.

Which statement is incorrect?

A.Every adapter must have a unique local MAC address.

B.A network adapter can be a member of only one group.

C.Solaris 9 requires test addresses for every IPMP group with more than one adapter.

D.Solaris 10 must have test addresses configured in an IPMP group with more than one adapter.

Answer: D

207. In Sun Cluster 3.2 a logical host's address is flagged as deprecated.

Which statement is true for the deprecated flag?

A.The deprecated flag is applicable for logical hosts and test addresses only.

B.The logical host address is always the source address in an outbound connection.

C.The logical address is never picked as the source address in an outbound connection.

D.It is impossible to remove the deprecated flag on a logical host by an additional resource.

Answer: C

208. Sun Cluster 3.2 requires an IPMP configuration for a logical host resource.

Which action does NOT lead to a persistent IPMP configuration?

A.Let the cluster scinstall command configure the IPMP group.

B.Modify the /etc/default/mpathd configuration file if needed.

C.Configure the IPMP group with ifconfig at the command prompt.

D.Configure the IPMP group in /etc/hostname. file manually.

Answer: C

209. Which statement is incorrect regarding the in.mpathd daemon?

A.IPMP uses ping to contact any router listed in the routing table.

B.If no router exists, IPMP submits a ping to the all-hosts multicast address.

C.If static host routes are available, these hosts will be contacted using ping.

D.You can configure link state monitoring on Solaris 9.

Answer: D

210. IPMP itself is cluster unaware, so Sun Cluster 3.2 must interact with the IPMP software.

Which activity is NOT performed by Sun Cluster 3.2?

A.The cluster wide IPMP status is stored in the CCR.

B.Sun Cluster 3.2 facilitates a failover to a different node if a complete IPMP group fails on a node.

C.Sun Cluster 3.2 facilitates an adapter failback if an unhealthy adapter was repaired.

D.Sun Cluster's pnmd monitors the IPMP status.

Answer: C

211. You have just configured an IPMP group on each cluster node, and want to test that they are running correctly before the cluster goes live. However, the server is physically located in a closed datacenter.

What should you do?

A.You CANNOT test the functionality of IPMP without physically pulling out a network cable.

B.Run the commands on each node:

if_mpadm -d

if_mpadm -r

C.Run the command on one node:

clrs disable -g lh-rg lh-res

in which lh-res refers to the logical IP address in the lh-rg resource group.

D.Run:

kill -HUP in.mpathd

Answer: B

212. What are the requirements if you want to use only IPV6 test address for your public network adapters, and no IPV4 test address?

A.You cannot use IPV6 for the IPMP interfaces.

B.You can use IPV6, but you must have at least one IPV4 test address.

C.You must not have IPV6 configured on any servers outside the cluster.

D.You must have at least one server outside the cluster configured using IPV6.

Answer: D

213. In Sun Cluster 3.2, what command is used to view the status of IPMP group members on all nodes?

A.scstat -IPMP

B.IPMP status -v

C.clnode status -m

D.clrg public -v all

Answer: C

214. Which IPMP restriction for multiple adapters in an IPMP group is no longer a restriction in Solaris 10?

A.Every adapter must have a dedicated test IP address.

B.A network adapter can be a member of only one group.

C.Each Ethernet adapter must have a unique MAC address.

D.All members of the same group must be on the same subnet.

Answer: A

215. Your network administrator has disabled all routers from answering pings. You notice that IPMP is reporting network path failures and this is causing your RGs to fail over.

Which of the following actions will fix the problem?

A.Kill the pnmd daemon.

B.Change FAILBACK to NO in /etc/default/mpathd file.

C.Increase the FAILURE_DETECTION_TIME in /etc/default/mpathd file.

D.Add static routes to IP addresses that answer pings on your subnet.

Answer: D

216. When configuring IPMP groups for use in the Sun Cluster 3.2 environment, which two statements are true? (Choose two.)

A.All interfaces must be in IPMP groups.

B.Each interface can be a member of multiple IPMP groups.

C.Adapters within the same IPMP group must be of the same transport type.

D.Adapters on different subnets must be in different IPMP groups.

Answer: CD

217. While setting up logical network interfaces within IPMP groups, what is the maximum number of virtual network interfaces for each physical adapter?

A.You may have only two virtual interfaces for each physical adapter.

B.You may have an unlimited number of virtual interfaces for each physical adapter.

C.You can have a maximum of 8192 virtual interfaces for each physical adapter.

D.You can have as many virtual interfaces as you have physical nodes in the cluster.

Answer: C

218. When you create a new resource group with a command such as

clrg create -n vincent,theo,noodle mynewrg

why would you be concerned about the order of the nodes that you list ?

A.The node list is always an ordered list.

B.The resource group will always fail back to the node that is listed first in the node list.

C.You have to list the storage-connected nodes before any other nodes not connected to the storage.

D.You need not be concerned with the order. The meaning is the same regardless of the order in which the nodes are listed.

Answer: A

219. You have an application resource web-res representing a web resource, and you set the property

Resource_dependencies_restart=ora-server-res to represent that the web resource is dependent on an Oracle resource.

Which two are true? (Choose two.)

A.The cluster RGM will try to start the web resource only if the Oracle resource is already running.

B.If the RGM restarts the web resource it will then also restart the Oracle resource.

C.If the RGM restarts the Oracle resource it will then also restart the web resource.

D.The Oracle resource must be configured in the same resource group as the web resource.

E.If the web resource fails over, the Oracle resource must be restarted.

Answer: AC

220. You need to have two failover IP addresses on separate subnets both fail over along with an application. How do you create your IP address resources?

A.You can create a single LogicalHostname resource respresenting both IP addresses.

B.You must create two separate LogicalHostname resources and put them in the same resource group as the application resource.

C.You must create two separate LogicalHostname resources and put them in two separate resource groups. A resource group can only contain one LogicalHostname resource.

D.You must create a SharedAddress resource. Only a SharedAddress resource can refer to multiple IP addresses.

Answer: B

221. The SUNW.HAStoragePlus resource manages different types of storage. In which case is it mandatory to configure HAStoragePlus?

A.to manage global raw devices

B.to manage global file systems

C.to manage failover file systems

D.to manage global file system devices with an Oracle RAC configuration

Answer: C

222. A new cluster has been installed and the customer needs to create a HA-NFS service. Which steps have to be executed ?

A.Register SUNW.nfs dataservice, create a resource group, and create a SUNW.nfs resource.

B.Register SUNW.nfs dataservice, create a resource group, create a SUNW.HAStoragePlus resource, and create NFS resource.

C.Register SUNW.nfs dataservice, create a resource group, create a SUNW.Logicalhostname resource, create a SUNW.HAStoragePlus resource and create an NFS resource.

D.Create a resource group, create a SUNW.Logicalhostname resource, create a SUNW.HAStoragePlus resource.

Answer: C

223. A customer needs to do some maintenance to an application resource in their NFS resource group. They want to have NFS resources active, but want to be sure that if anything goes wrong during the operation, the NFS service is NOT restarted or the resource group switched.

How can this be accomplished ?

A.The customer can suspend the resource group.

B.The customer can place the resource group into a quiescent state.

C.The customer can place the resource group offline and then enable only the required resources.

D.The customer can keep the resource group online and then disable monitoring for all the resources.

Answer: A

224. A customer has a resource group called nfs-rg with a SUNW.nfs resource called nfs-rs. They want to check how the Failover_mode property has been set.

Which is the correct command used to accomplish this task ?

A.clrg status

B.clrs show -g nfs-rg

C.clrs show -v nfs-rs

D.clrs list -v nfs-rs

Answer: C

225. When you want to fail over your application between zones, you have to understand the principles of integrating non-global zones in a cluster.

Which statement below is incorrect for the zones integration?

A.An application can fail over between 2 zones on the same node.

B.The zone names for one resource group can be different on the various nodes.

C.All the zones have to be booted with the HA Container agent.

D.You do NOT have to specify an IP address for each zone within zonecfg.

Answer: C

226. Every resource in Sun Cluster 3.2 has to have a resource type. Some of the resource types are pre-registered, and most of them you have to register.

Which resource type listed below is pre-registered?

A.SUNW.gds

B.SUNW.nfs

C.SUNW.LogicalHostname

D.SUNW.HAStoragePlus

Answer: C

227. Which task below is correct when creating a storage resource for a failover file system?

A.You must add the entries for zfs file systems to /etc/vfstab.

B.You have to configure the mount at boot option to yes in /etc/vfstab.

C.You have to add an entry for /etc/vfstab on each node for ufs and vxfs file systems.

D.If your file system is on a vxvm disk group, you have to import the disk group manually for the first time.

E.You have to enable your storage resource created with clrs create in a resource group where resources are already online.

Answer: C

228. Which components need to be available before creating an NFS resource called nfs-rs?

A.A /global/zfs/admin directory and dfstab.nfs-zs file.

B.On each node a SUNW.nfs directory and a dfstab.nfs-rs file are required in local storage.

C.A SUNW.nfs directory and a dfstab.nfs-rs file located on a global or failover file system.

D.Nothing has to be created, but the components must be available before the resource is enabled.

Answer: C

229. You want to suspend a resource group to perform application maintenance.

Which statement is incorrect regarding this task?

A.The state of an application resource is reported as offline if the application administrator stops the application manually.

B.You can still enable or disable a resource with the clrs command while the resource group is suspended.

C.When the node which masters the resource group fails, the resource group will not be failed over to another node.

D.If the entire cluster reboots, the suspended resource group will remain offline.

Answer: A

230. You have a three node cluster with three different resource groups. Rg-1 can only run on node1 and node3, Rg-2 can only run on node2 and node3, while Rg-3 can only run on node3. If Rg-1 or Rg-2 move to node3, Rg-3 has to be stopped.

How would you accomplish this ?

A.clrs set -p RG_affinities=--Rg-1,--Rg-2 Rg-3

B.clresourcegroup set -p RG_affinities=--Rg-3 +

C.clresourcegroup set -c RG_affinities=--Rg-1,Rg-2 Rg-3

D.clresourcegroup set -p RG_affinities=--Rg-1,--Rg-2 Rg-3

Answer: D

231. Suppose you have a resource apache-rs and you have been asked to make sure that this resource can be restarted on the same node, but must NOT trigger a failover to another node.

Which command will accomplish this task?.

A.clrs set -p Log_only=false apache-rs

B.clrs set -p Restart_only=true apache-rs

C.clrs set -p Failover_mode=RESTART_ONLY apache-rs

D.clrs set -p Failover_enabled=RESTART_ONLY apache-rs

Answer: C

232. Suppose you have a resource apache-rs and you were informed that its state is in STOP_FAILED on node1.

Which command is needed prior to starting this resource again?

A.clrs set -n node1 -f STOP_FAILED apache-rs

B.clrs clear -n node1 -f STOP_FAILED apache-rs

C.clrs reset -n node1 -f STOP_FAILED apache-rs

D.clrs disable -n node1 -f STOP_FAILED apache-rs

Answer: B

233. Suppose you have two nodes called node1 and node2, and each of them has a non-global zone running called zone1.

Which two commands register a failover resource group named apache-rg to be potentially mastered on both non-global zones? (Choose two.)

A.clrg create -z zone1 apache-rg

B.clrg create -n node1,node2 -z zone1 apache-rg

C.clrg create -n zone1@node1,zone1@node2 apache-rg

D.clrg create -n node1:zone1,node2:zone1 apache-rg

Answer: BD

234. A cluster administrator should be aware of the capabilities of weak and strong affinities.

Which statement is incorrect?

A.Weak affinities are denoted by one single + or - sign.

B.Weak affinities can be overruled by manual switchovers.

C.Weak and strong affinities can be applied together on one resource group.

D.If you defined a weak negative affinity between rg1 and rg2, and this affinity CANNOT be satisfied, rg1 will NOT fail over on a node failure.

Answer: D

235. Which should be taken into consideration when configuring a scalable service?

A.Data must be on global file system, file locking is managed from PXFS.

B.Data can be on a local highly available file system and accessed using cluster interconnect.

C.Data must be on a global file system, file locking must be managed from Application.

D.Data can be on a global file system, file locking is managed from PXFS.

Answer: C

236. You have two applications. You prefer they run on different nodes because they both use a lot of memory, but

you want to preserve high availability for both applications even if only one node is left alive.

Which kind of resource group affinity should you set between the two resource groups?

A.weak positive affinity

B.weak negative affinity

C.strong negative affinity

D.strong positive affinity with failover delegation

Answer: B

237. Which two of the following properties, if set to a value greater than one, will automatically identify a resource group as scalable? (Choose two.)

A.Nodelist

B.Scalable_count

C.Desired_primaries

D.Maximum_primaries

Answer: CD

238. The cluster administrator wants to use your cluster for test and production systems in an active mode, with resource groups called test-rg and prod-rg. The cluster administrator's goal is that test-rg and prod-rg will never run on the same node. Whenever there is only one node left for prod-rg, and test-rg is running on that node, test-rg has to be switched off.

Which command will achieve this behavior?

A.clrg set -p RG_affinities=-test-rg prod-rg

B.clrg set -p RG_affinities=-prod-rg test-rg

C.clrg set -p RG_affinities=--test-rg prod-rg

D.clrg set -p RG_affinities=--prod-rg test-rg

Answer: D

239. RG1 is a scalable group and RG2 is a failover group. Which two commands will succeed? (Choose two.)

A.clrg set -p RG_affinities=++RG2 RG1

B.clrg set -p RG_affinities=++RG1 RG2

C.clrg set -p RG_affinities=--RG1 RG2

D.clrg set -p RG_affinities=+++RG2 RG1

Answer: BC

240. Which two are true about Oracle10g RAC database instances in Sun Cluster 3.2? (Choose two.)

A.The Oracle Cluster Ready Services (CRS) controls starts and stops of the RAC instances.

B.In the Sun Cluster 3.2 environment, you disable Cluster Ready Services (CRS) and let the cluster resources start and stop the instances.

C.You can create Sun Cluster 3.2 resources that can act as a proxy to Cluster Ready Services (CRS) for starting and stopping the RAC instances.

D.You must create separate Sun Cluster 3.2 resources in separate resource groups for controlling the RAC database instance on each node.

Answer: AC

241. You configure Oracle10g RAC in Sun Cluster 3.2 with the data on a global file system. What will happen?

A.You will get an error from the database configuration utility.

B.It will actually configure and run, but it is not a supported Oracle configuration.

C.It will run correctly at first, but will fail if the entire cluster ever reboots.

D.It will seem to configure correctly. The database software will actually check the configuration and return an error at run time.

Answer: B

242. After consolidating multiple clusters into a single cluster you will need to run multiple HA Oracle instances in the same cluster.

Which of the steps is unnecessary in an environment where the Oracle instances have to be accessed over a network?

A.Configure listener resources for the Oracle instances.

B.Configure the server resources for the Oracle instances.

C.Configure your global/failover file system in HAStoragePlus resources.

D.Configure the resources for Oracle Grid Control to administer the Oracle instances.

Answer: D

243. Click the Task button.

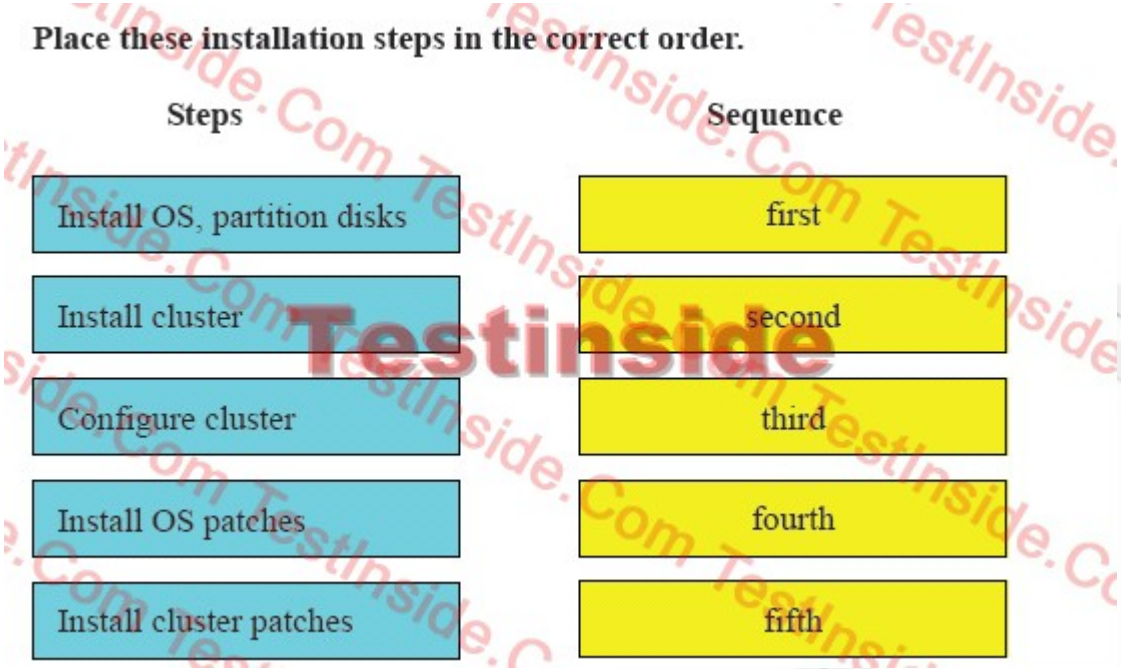

Answer: